Reviewer 1

Comments

Reviewer 1 - Initial submission, 23/10/2026
This paper proposes AMORE, a multimodal recommendation framework that connects edge-level affinity modulation (ELAM) with affinity-guided coarse-to-fine decoupling (CFDM) and modality-adaptive contrastive alignment (MACA). The problem is meaningful and the paper is well written, but the improvements are marginal, the ablations are unconvincing, and the comparison and evaluation are incomplete. I do not think it meets the bar of IP&M in its current form.
Strengths:
1. Carrying one interaction-derived affinity signal from graph modulation into feature fusion, while keeping the ID pathway intact, is a coherent and well-motivated design idea.
2. The paper is well organized, and the revision is responsive: explicit research objectives and neatly grouped result tables.
3. The evaluation covers twelve baselines on three datasets, plus ablations, sensitivity analysis, and visualizations, giving a reasonably complete empirical picture.
Weaknesses:
1. The gains over the strongest baseline (MENTOR) are marginal and inconsistent. Some of the twelve entries in Table 2 favor MENTOR (e.g., Recall@20 on Baby: 0.1040 vs. 0.1048), and most others are around or below 1%. All results are single point estimates without variance or significance tests.
2. The ablation study does not support the flagship module. Removing ELAM changes Recall@20 by only 0.0002 on Baby and 0.0005 on Sports, which may well fall within training noise.
3. Several closely related methods discussed at length, such as EgoGCN, LGMRec, EVEN, MMSSL and DiffMM, are not included in the comparison.
4. All datasets are Amazon product categories from a single source, so the conclusions may not transfer to other domains such as short-video or music recommendation.
5. No efficiency analysis is reported; training/inference time and memory should be compared against the baselines.

Reviewer 2

Comments

Reviewer 2 - Initial submission, 16/10/2026
1. The central claimed contribution, namely the use of interaction-derived affinity as the connecting signal between graph modulation and feature fusion, is never directly validated: the ablations remove entire modules and therefore cannot isolate this mechanism.
2. The pseudocode in Algorithm 1 does not match the mathematical formulation in the body. Every symbol should be defined and consistent with Equations.
3. Baseline comparability: most scores are quoted, not rerun under a unified protocol. I strongly recommend comparing under the same experimental conditions. The current experiment is not convincing and the gains also are marginal.
4. The logicality of the manuscript's writing is insufficient, cold-start and long-tail are motivated in the introduction, but never evaluated.

5. Internal inconsistencies and typos need to be proofread, such as the numberof baselines, hyperparameter typos and etc.
In its current form, however, I recommend rejection.

Reviewer 3

Comments

Reviewer 3 - Initial submission, 31/10/2026
This paper proposes AMORE, a multimodal recommendation framework built upon edge level affinity modulation and coarse to fine feature decoupling. The overall technical idea is reasonable and the framework is complete. However, several critical weaknesses remain, including insufficient literature coverage, inadequate theoretical interpretation of core loss functions, incomplete experimental validations, and missing key supplementary experiments corresponding to the claimed denoising capability.

Review Comments ：
1. The literature coverage in the introduction is insufficient, and recent achievements on multimodal graph denoising recommendation are missing.
2. The theoretical value of the joint loss in MACA is inadequately analyzed. It does not explain the gain brought by the mean variance alignment loss compared with InfoNCE alone, although MACA serves as one of the core modules of this work.
3. Experimental results lack repeated tests with multiple random seeds, standard deviations and significance tests. Minor performance improvements cannot rule out the possibility of random fluctuations.
4. Current ablation experiments are relatively coarse grained. Additional group wise validation is suggested to verify the independent contribution when the affinity signal is only used for the graph modulation branch without guiding the decoupling module, as well as the effect of removing specific auxiliary losses.
5. The existing parameter sensitivity experiments only present the sensitivity of the self supervised contrastive learning weight. Sensitivity analyses for the weight of the Gaussian distribution alignment loss and the L2 regularization term are absent.
6. Controlled noise robustness experiments are missing. Though the paper claims to dynamically filter noisy interactions, comparisons with relevant baselines under artificially injected interaction or modal noise of different ratios have not been conducted, even if the authors state that controlled noise tests will be left for future work.
7. Model efficiency comparisons are lacking. Metrics such as parameter quantity, per epoch training time and GPU memory consumption are not reported, making it impossible to evaluate the computational overhead introduced by the newly added modules.
8. In the analysis of certain experimental sections, authors should focus on analyzing the internal mechanism of proven models or modules, and avoid analyzing conclusions that have not been validated in the paper.